

# Supplementary Information for Improving Probabilistic Models in Text Classification via Active Learning<sup>\*</sup>

Mitchell Bosley<sup>†‡</sup>      Saki Kuzushima<sup>†§</sup>      Ted Enamorado<sup>¶</sup>  
Yuki Shiraito<sup>||</sup>

First draft: September 10, 2020

This draft: September 23, 2022

## Contents

<b>A</b>	<b>Detailed explanations about the EM algorithm to estimate parameters</b>	<b>1</b>
<b>B</b>	<b>EM algorithm for binary classification with multiple clusters</b>	<b>4</b>
B.1	Summary . . . . .	4
B.2	Model . . . . .	4
B.3	Results . . . . .	6
<b>C</b>	<b>Multiclass Classification</b>	<b>8</b>
C.1	Model . . . . .	8
C.2	Results . . . . .	9

---

<sup>\*</sup>We thank Ken Benoit, Yaoyao Dai, Chris Fariss, Yusaku Horiuchi, Kosuke Imai, Walter Mebane, Daichi Mochihashi, Kevin Quinn, and audiences at the 2020 Annual Meeting of the American Political Science Association, the 2021 Annual Meeting of the Midwest Political Science Association, the 11th Annual Conference on New Directions in Analyzing Text as Data, and the 2022 Summer Meeting of the Japanese Society for Quantitative Political Science, and seminar participants at the University of Michigan and members of the Junior Faculty Workshop at Washington University in St. Louis for useful comments and suggestions.

<sup>†</sup>These authors have contributed equally to this work.

<sup>‡</sup>Ph.D. Candidate, Department of Political Science, University of Michigan. Email: [mcbosley@umich.edu](mailto:mcbosley@umich.edu).

<sup>§</sup>Ph.D. Candidate, Department of Political Science, University of Michigan. Email: [skuzushi@umich.edu](mailto:skuzushi@umich.edu)

<sup>¶</sup>Assistant Professor, Department of Political Science, Washington University in St. Louis. Siegle Hall, 244. One Brookings Dr. St Louis, MO 63130-4899. Phone: 314-935-5810, Email: [ted@wustl.edu](mailto:ted@wustl.edu), URL: [www.tedenamorado.com](http://www.tedenamorado.com).

<sup>||</sup>Assistant Professor, Department of Political Science, University of Michigan. Center for Political Studies, 4259 Institute for Social Research, 426 Thompson Street, Ann Arbor, MI 48104-2321. Phone: 734-615-5165, Email: [shiraito@umich.edu](mailto:shiraito@umich.edu), URL: [shiraito.github.io](http://shiraito.github.io).

<b>D</b>	<b>Model Specifications and Description of the Datasets in the Validation Performance</b>	<b>12</b>
D.1	Pre-processing . . . . .	12
D.2	Datasets . . . . .	12
<b>E</b>	<b>Additional Results on Classification Performance</b>	<b>14</b>
<b>F</b>	<b>Main Results when Varying Positive Class Rate</b>	<b>16</b>
<b>G</b>	<b>Visual Demonstration of Active Keyword</b>	<b>17</b>
<b>H</b>	<b>Classification Performance with Mislabels</b>	<b>19</b>
H.1	Mislabeled Keywords . . . . .	19
H.2	Mislabeled Documents . . . . .	21
<b>I</b>	<b>Comparison of the predictions between <i>activeText</i> and xgboost predictions for the Gohdes (2020) data</b>	<b>23</b>
<b>J</b>	<b>Regression Table in Gohdes (2020)</b>	<b>25</b>
<b>K</b>	<b>Effect of Labeling More Sentences for the Park et al. (2020) Reanalysis</b>	<b>29</b>

Note that to facilitate exposition, in the main text, we use the words political and non-political labels to describe the problem of binary classification. Without loss of generality, in this supplemental information material, we use the positive vs. negative class dichotomy instead.

## A Detailed explanations about the EM algorithm to estimate parameters

Let  $\mathbf{D}^{lp}$ ,  $\mathbf{D}^{ln}$  and  $\mathbf{D}^u$  be the document feature matrices for documents with positive labels, documents with negative labels, and unlabeled documents, respectively. Also let  $N^{lp}$ ,  $N^{ln}$ , and  $N^u$  be the number of documents with positive labels, negative labels, and documents without labels. Likewise,  $\mathbf{C}^{lp}$  and  $\mathbf{C}^{ln}$  be the vectors of positive and negative labels. Then, the observed-data likelihood is:

$$\begin{aligned}
& p(\pi, \boldsymbol{\eta} | \mathbf{D}, \mathbf{C}^{lp}, \mathbf{C}^{ln}) \\
& \propto p(\pi) p(\boldsymbol{\eta}) p(\mathbf{D}^{lp}, \mathbf{C}^{lp} | \pi, \boldsymbol{\eta}) p(\mathbf{D}^{ln}, \mathbf{C}^{ln} | \pi, \boldsymbol{\eta}) \left[ p(\mathbf{D}^u | \pi, \boldsymbol{\eta}) \right]^\lambda \\
& = p(\pi) p(\boldsymbol{\eta}) \times \prod_{i=1}^{N^{lp}} p(\mathbf{D}_i^{lp} | Z_i = 1, \boldsymbol{\eta}) p(Z_i = 1 | \pi) \times \prod_{i=1}^{N^{ln}} \left\{ p(\mathbf{D}_i^{ln} | Z_i = 0, \boldsymbol{\eta}) p(Z_i = 0 | \pi) \right\} \\
& \quad \times \left[ \prod_{i=1}^{N^u} \left\{ p(\mathbf{D}_i^u | Z_i = 1, \boldsymbol{\eta}) p(Z_i = 1 | \pi) + p(\mathbf{D}_i^u | Z_i = 0, \boldsymbol{\eta}) p(Z_i = 0 | \pi) \right\} \right]^\lambda \tag{1} \\
& \propto \underbrace{\left\{ (1 - \pi)^{\alpha_0 - 1} \prod_{v=1}^V \eta_{v0}^{\beta_{0v} - 1} \right\} \times \left\{ \pi^{\alpha_1 - 1} \prod_{v=1}^V \eta_{v1}^{\beta_{1v} - 1} \right\}}_{\text{prior}} \times \underbrace{\prod_{i=1}^{N^{lp}} \left\{ \prod_{v=1}^V \eta_{v1}^{D_{iv}} \times \pi \right\}}_{\text{positive labeled doc. likelihood}} \\
& \quad \times \underbrace{\prod_{i=1}^{N^{ln}} \left\{ \prod_{v=1}^V \eta_{v0}^{D_{iv}} \times (1 - \pi) \right\}}_{\text{negative labeled doc. likelihood}} \times \underbrace{\left[ \prod_{i=1}^{N^u} \left\{ \prod_{v=1}^V \eta_{v0}^{D_{iv}} \times (1 - \pi) \right\} + \left\{ \prod_{v=1}^V \eta_{v1}^{D_{iv}} \times \pi \right\} \right]^\lambda}_{\text{unlabeled doc. likelihood}}
\end{aligned}$$

We weigh the part of the observed likelihood that refers to the unlabeled document with  $\lambda \in [0, 1]$ . This is done because we typically have many more unlabeled documents than labeled documents. By downweighting the information from the unlabeled document (i.e., setting  $\lambda$  to be small), we can use more reliable information from labeled documents than from unlabeled documents.

We estimate the parameters  $\pi$  and  $\boldsymbol{\eta}$  using EM algorithm Dempster et al. (1977) and our implementation is presented as pseudocode in Algorithm 1. Note that by taking the

---

**Algorithm 1:** EM algorithm to classify text

---

**Result:** Maximize  $p(\pi^{(t)}, \boldsymbol{\eta}^{(t)} \mid \mathbf{D}^l, \mathbf{Z}^l, \mathbf{D}^u, \boldsymbol{\alpha}, \boldsymbol{\beta})$   
**if** *In the first iteration of Active learning* **then**  
    Initialize  $\pi$  and  $\boldsymbol{\eta}$  by Naive Bayes;  
     $\pi^{(0)} \leftarrow \text{NB}(\mathbf{D}^l, \mathbf{Z}^l, \boldsymbol{\alpha})$ ;  
     $\boldsymbol{\eta}^{(0)} \leftarrow \text{NB}(\mathbf{D}^l, \mathbf{Z}^l, \boldsymbol{\beta})$ ;  
**else**  
    Inherit  $\pi^{(0)}$  and  $\boldsymbol{\eta}^{(0)}$  from the previous iteration of Active learning;  
**end**  
**while**  $p(\pi^{(t)}, \boldsymbol{\eta}^{(t)} \mid \mathbf{D}^l, \mathbf{Z}^l, \mathbf{D}^u, \boldsymbol{\alpha}, \boldsymbol{\beta})$  *does not converge* **do**  
    (1) E step: obtain the probability of the class for unlabeled documents;  
     $p(\mathbf{Z}^u \mid \pi^{(t)}, \boldsymbol{\eta}^{(t)} \mathbf{D}^l, \mathbf{Z}^l, \mathbf{D}^u) \leftarrow \text{E step}(\mathbf{D}^u, \pi^{(t)}, \boldsymbol{\eta}^{(t)})$ ;  
    (2) Combine the estimated classes for the unlabeled docs and the known classes  
    for the labeled docs;  
     $p(\mathbf{Z} \mid \pi^{(t)}, \boldsymbol{\eta}^{(t)}, \mathbf{D}^l, \mathbf{Z}^l, \mathbf{D}^u) \leftarrow \text{combine}(\mathbf{D}^l, \mathbf{D}^u, \mathbf{Z}^l, p(\mathbf{Z}^u \mid \pi^{(t)}, \boldsymbol{\eta}^{(t)}, \mathbf{D}^l, \mathbf{Z}^l, \mathbf{D}^u))$ ;  
    (3) M step: Maximize  $Q \equiv \mathbb{E}[p(\pi, \boldsymbol{\eta}, \mathbf{Z}^u \mid \mathbf{D}^l, \mathbf{Z}^l, \mathbf{D}^u, \boldsymbol{\alpha}, \boldsymbol{\beta})]$  w.r.t  $\pi$  and  $\boldsymbol{\eta}$ ;  
     $\pi^{(t+1)} \leftarrow \arg\max Q$ ;  
     $\boldsymbol{\eta}^{(t+1)} \leftarrow \arg\max Q$ ;  
    (4) Check convergence: Obtain the value of  $p(\pi^{(t+1)}, \boldsymbol{\eta}^{(t+1)} \mid \mathbf{D}^l, \mathbf{Z}^l, \mathbf{D}^u, \boldsymbol{\alpha}, \boldsymbol{\beta})$ ;  
**end**

---

expectation of the log complete likelihood function (Q function),

$$\begin{aligned}
Q &\equiv \mathbb{E}_{\mathbf{Z} \mid \pi^{(t)}, \boldsymbol{\eta}^{(t)}, \mathbf{D}, \mathbf{C}}[p(\pi, \boldsymbol{\eta}, \mathbf{Z} \mid \mathbf{D}, \mathbf{C})] \\
&= (\alpha_0 - 1) \log(1 - \pi^{(t)}) + (\alpha_1 - 1) \log \pi^{(t)} + \sum_{v=1}^V \left\{ (\beta_{0v} - 1) \log \eta_{v0}^{(t)} + (\beta_{1v} - 1) \log \eta_{v1}^{(t)} \right\} \\
&\quad + \sum_{i=1}^{N^{lp}} \left\{ \sum_{v=1}^V D_{iv} \log \eta_{v1}^{(t)} + \log \pi^{(t)} \right\} + \sum_{i=1}^{N^{ln}} \left\{ \sum_{v=1}^V D_{iv} \log \eta_{v0}^{(t)} + \log(1 - \pi^{(t)}) \right\} \\
&\quad + \lambda \left[ \sum_{i=1}^{N^u} p_{i0} \left\{ \sum_{v=1}^V D_{iv} \log \eta_{v0}^{(t)} + \log(1 - \pi^{(t)}) \right\} + p_{i1} \left\{ \sum_{v=1}^V D_{iv} \log \eta_{v1}^{(t)} + \log \pi^{(t)} \right\} \right]
\end{aligned} \tag{2}$$

where  $p_{ik}$  is the posterior probability of a document  $i$  being assigned to the  $k$  th cluster,  $k = \{0, 1\}$ , given data and the parameters at  $t$  th iteration. If a document has a positive label,  $p_{i0} = 0$  and  $p_{i1} = 1$ .

If a document has no label,

$$p_{i0} = 1 - p_{i1}$$

$$p_{i1} = \frac{\prod_{v=1}^V \eta_{v1}^{D_{iv}} \times \pi}{\prod_{v=1}^V \{\eta_{v0}^{D_{iv}} \times (1 - \pi)\} + \prod_{v=1}^V \{\eta_{v1}^{D_{iv}} \times \pi\}} \quad (3)$$

Equation 3 also works as the prediction equation. The predicted class of a document  $i$  is  $k$  that maximizes this posterior probability.

In the M-step, we maximize the Q function, and obtain the updating equations for  $\pi$  and  $\eta$ . The updating equation for  $\pi$  is the following.

$$\pi^{(t+1)} = \frac{\alpha_1 - 1 + N^{lp} + \lambda \sum_{i=1}^{N^u} p_{i1}}{\left(\alpha_1 - 1 + N^{lp} + \lambda \sum_{i=1}^{N^u} p_{i1}\right) + \left(\alpha_0 - 1 + N^{ln} + \lambda \sum_{i=1}^{N^u} p_{i0}\right)} \quad (4)$$

The updating equation for  $\eta$  is the following.

$$\begin{aligned} \hat{\eta}_{v0}^{(t+1)} &\propto (\beta_{v0} - 1) + \sum_{i=1}^{N^{ln}} D_{iv} + \lambda \sum_{i=1}^{N^u} p_{i0} D_{iv}, \quad v = 1, \dots, V \\ \hat{\eta}_{v1}^{(t+1)} &\propto (\beta_{v1} - 1) + \sum_{i=1}^{N^{lp}} D_{iv} + \lambda \sum_{i=1}^{N^u} p_{i1} D_{iv}, \quad v = 1, \dots, V \end{aligned} \quad (5)$$

## B EM algorithm for binary classification with multiple clusters

### B.1 Summary

The model outlined above assumes that there are two latent clusters, each linked to the positive and the negative class. However, this assumption can be relaxed to link multiple clusters to the negative class. In the world of mixture models, the simplest setup is to let  $K = 2$  since the classification goal is binary, and we can link each latent cluster to the final classification categories. A more general setup is to use  $K > 2$  even when a goal is a binary classification. If  $K > 2$ , but our focus is to uncover the identity of one cluster, we can choose one of the latent clusters to be linked to the “positive” class and let all other latent clusters be linked to the “negative” class (see e.g., Larsen and Rubin 2001 for a similar idea in the realm of record linkage). In other words, we collapse the  $K - 1$  latent clusters into one class for the classification purpose. Using  $K > 2$  makes sense if the “negative” class consists of multiple sub-categories. For instance, suppose researchers are interested in classifying news articles into political news or not. Then, it is reasonable to assume that the non-political news category consists of multiple sub-categories, such as technology, entertainment, and sports news.

### B.2 Model

This section presents a model and inference algorithm when we use more than 2 latent clusters in estimation but the final classification task is binary. In other words, we impose a hierarchy where many latent clusters are collapsed into the negative class. In contrast, the positive class is made out of just one class. The model presented is as follows:

$$\begin{aligned}
 \pi &\sim \text{Dirichlet}(\boldsymbol{\alpha}) \\
 Z_i &\stackrel{i.i.d}{\sim} \text{Categorical}(\boldsymbol{\pi}) \\
 \eta_k &\stackrel{i.i.d}{\sim} \text{Dirichlet}(\boldsymbol{\beta}_k), \quad k = \{1, \dots, K\} \\
 \mathbf{D}_i | Z_i = k &\stackrel{i.i.d}{\sim} \text{Multinomial}(n_i, \boldsymbol{\eta}_k)
 \end{aligned} \tag{6}$$

Note that  $\boldsymbol{\pi}$  is now a probability vector of length  $K$ , and it is drawn from a Dirichlet distribution.

Let  $k^*$  be the index of the cluster linked to the positive class. The observed likelihood is

the following.

$$\begin{aligned}
& p(\boldsymbol{\pi}, \boldsymbol{\eta} | \mathbf{D}, \mathbf{C}^{lp}, \mathbf{C}^{ln}) \\
& \propto p(\boldsymbol{\pi}) p(\boldsymbol{\eta}) p(\mathbf{D}^{lp}, \mathbf{C}^{lp} | \boldsymbol{\pi}, \boldsymbol{\eta}) p(\mathbf{D}^{ln}, \mathbf{C}^{ln} | \boldsymbol{\pi}, \boldsymbol{\eta}) \left[ p(\mathbf{D}^u | \boldsymbol{\pi}, \boldsymbol{\eta}) \right]^\lambda \\
& = p(\boldsymbol{\pi}) p(\boldsymbol{\eta}) \times \prod_{i=1}^{N^{lp}} p(\mathbf{D}_i^{lp} | Z_i = k^*, \boldsymbol{\eta}) p(Z_i = k^* | \boldsymbol{\pi}) \\
& \quad \times \prod_{i=1}^{N^{ln}} \sum_{k \neq k^*} \left\{ p(\mathbf{D}_i^{ln} | Z_i = k, \boldsymbol{\eta}) p(Z_i = k | \boldsymbol{\pi}) \right\} \times \left[ \prod_{i=1}^{N^u} \sum_{k=1}^K \left\{ p(\mathbf{D}_i^u | Z_i = k, \boldsymbol{\eta}) p(Z_i = k | \boldsymbol{\pi}) \right\} \right]^\lambda \\
& \propto \underbrace{\prod_{k=1}^K \left\{ \pi_k^{\alpha_k - 1} \prod_{v=1}^V \eta_{vk}^{\beta_{kv} - 1} \right\}}_{\text{prior}} \times \underbrace{\prod_{i=1}^{N^{lp}} \left\{ \prod_{v=1}^V \eta_{vk^*}^{D_{iv}} \times \pi_{k^*} \right\}}_{\text{positive labeled doc. likelihood}} \\
& \quad \times \underbrace{\prod_{i=1}^{N^{ln}} \sum_{k \neq k^*} \left\{ \prod_{v=1}^V \eta_{vk}^{D_{iv}} \times \pi_k \right\}}_{\text{negative labeled doc. likelihood}} \times \underbrace{\left[ \prod_{i=1}^{N^u} \sum_{k=1}^K \left\{ \prod_{v=1}^V \eta_{vk}^{D_{iv}} \times \pi_k \right\} \right]^\lambda}_{\text{unlabeled doc. likelihood}}
\end{aligned} \tag{7}$$

The Q function (the expectation of the complete log likelihood) is

$$\begin{aligned}
Q & \equiv \mathbb{E}_{\mathbf{Z} | \boldsymbol{\pi}^{(t)}, \boldsymbol{\eta}^{(t)}, \mathbf{D}, \mathbf{C}} [p(\boldsymbol{\pi}, \boldsymbol{\eta}, \mathbf{Z} | \mathbf{D}, \mathbf{C})] \\
& = \sum_{k=1}^K \left[ (\alpha_k - 1) \log \pi_k^{(t)} + \sum_{v=1}^V \left\{ (\beta_{kv} - 1) \log \eta_{vk}^{(t)} \right\} \right] \\
& \quad + \sum_{i=1}^{N^{lp}} \left\{ \sum_{v=1}^V D_{iv} \log \eta_{vk^*}^{(t)} + \log \pi_{k^*}^{(t)} \right\} + \sum_{i=1}^{N^{ln}} \sum_{k \neq k^*} p_{ik} \left\{ \sum_{v=1}^V D_{iv} \log \eta_{vk}^{(t)} + \log \pi_k^{(t)} \right\} \\
& \quad + \lambda \left[ \sum_{i=1}^{N^u} \sum_{k=1}^K p_{ik} \left\{ \sum_{v=1}^V D_{iv} \log \eta_{vk}^{(t)} + \log \pi_k^{(t)} \right\} \right]
\end{aligned} \tag{8}$$

The posterior probability of  $Z_i = k$ ,  $p_{ik}$ , is

$$p_{ik} = \frac{\prod_{v=1}^V \eta_{vk}^{D_{iv}} \times \pi_k}{\sum_{k=1}^K \left[ \prod_{v=1}^V \eta_{vk}^{D_{iv}} \times \pi_k \right]} \tag{9}$$




Figure B.1: **Classification Results with 2 and 5 Clusters.** The darker lines show the results with 5 latent clusters and the lighter lines show 2 latent clusters. The columns correspond to various proportions of positive labels in the corpus. The y-axis indicates the out-of-sample F1 score and the x-axis show the number of sampling steps. Using multiple clusters improves the classification performance when the number of latent clusters matches the data generating process.

M step estimators are The updating equation for  $\pi$  is the following.

$$\hat{\pi}_k \propto \begin{cases} \alpha_k - 1 + \sum_{i=1}^{N^{ln}} p_{ik} + \lambda \sum_{i=1}^{N^u} p_{ik} & \text{if } k \neq k^* \\ \alpha_k - 1 + N^{lp} + \lambda \sum_{i=1}^{N^u} p_{ik^*} & \text{if } k = k^* \end{cases} \quad (10)$$

The updating equation for  $\eta$  is the following.

$$\hat{\eta}_{vk} \propto \begin{cases} (\beta_k - 1) + \sum_{i=1}^{N^{ln}} p_{ik} D_{iv} + \lambda \sum_{i=1}^{N^u} p_{ik} D_{iv} & \text{if } k \neq k^* \\ (\beta_k - 1) + \sum_{i=1}^{N^{lp}} D_{iv} + \lambda \sum_{i=1}^{N^u} p_{ik^*} D_{iv} & \text{if } k = k^* \end{cases} \quad (11)$$

Note that we downweight the information from the unlabeled documents by  $\lambda$ , to utilize more reliable information from labeled documents.

### B.3 Results

Figure B.1 shows the results of a model with just two latent clusters vs. a model with 5 latent clusters but only two final classes (positive vs. negative). The darker lines show the results with 5 latent clusters and the lighter lines show the results with 2 latent clusters. Overall, the model with 5 clusters performs better or as well as the model with 2 clusters. The gain from using 5 clusters is the highest when the proportion of positive labels is small and when the size of labeled data is small.

Figure B.2 shows the results when the multiple cluster approach and keyword upweighting approaches are combined.



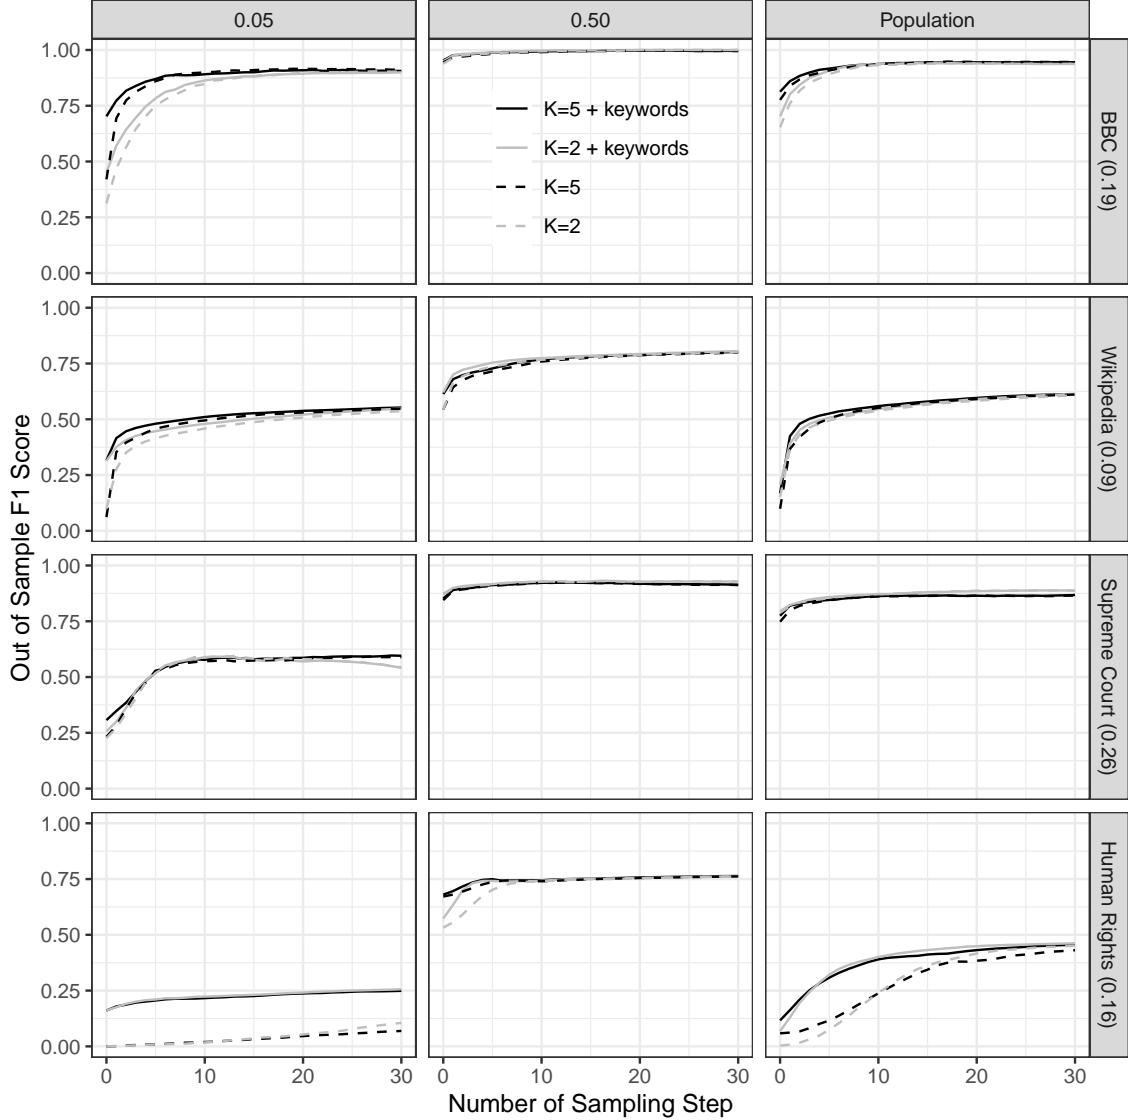


Figure B.2: **Classification Results with Multiple Clusters and Keywords.** The rows correspond to different datasets and the columns correspond to various proportions of positively labeled documents in the corpus. The y-axis indicates the out-of-sample F1 score and the x-axis show the number of sampling steps. The linetype show whether keywords are supplied: the solid lines show the results with keywords and the dashed lines without keywords. The colors show the number of latent clusters in the mixture model: the darker lines show the results with 5 latent clusters and the lighter lines with 2 latent clusters. Using 5 clusters leads to as good or slightly better performance than using 2 clusters. The performance improvement is the largest with the BBC corpus, which consists of 5 news topic categories. Likewise, our mixture models with keywords leads to as good or better performance than the models without keywords. The improvement is the largest with the human rights corpus, where the number of words per document is the smallest.

## C Multiclass Classification

### C.1 Model

This section presents a model and inference algorithm for multiclass classification. Let  $K$  be the number of the clusters and is equal to the number of classes to be classified, with  $K \geq 2$ . Differently than in SI B, we do not impose any hierarchies and the model is a true multi-class mixture model, where the end goal is to classify documents in  $K \geq 2$  classes. In other words, the model presented below is a generalization of the model presented in the main text.

$$\begin{aligned}
\pi &\sim \text{Dirichlet}(\boldsymbol{\alpha}) \\
Z_i &\stackrel{i.i.d}{\sim} \text{Categorical}(\boldsymbol{\pi}) \\
\eta_k &\stackrel{i.i.d}{\sim} \text{Dirichlet}(\boldsymbol{\beta}_k), \quad k = \{1, \dots, K\} \\
\mathbf{D}_i | Z_i = k &\stackrel{i.i.d}{\sim} \text{Multinomial}(n_i, \boldsymbol{\eta}_k)
\end{aligned} \tag{12}$$

Note that  $\boldsymbol{\pi}$  is now a probability vector of length  $K$ , and it is drawn from a Dirichlet distribution.

The observed likelihood is the following.

$$\begin{aligned}
p(\boldsymbol{\pi}, \boldsymbol{\eta} | \mathbf{D}, \mathbf{C}^l) &\propto p(\boldsymbol{\pi}) p(\boldsymbol{\eta}) p(\mathbf{D}, \mathbf{C} | \boldsymbol{\pi}, \boldsymbol{\eta}) \left[ p(\mathbf{D}^u | \boldsymbol{\pi}, \boldsymbol{\eta}) \right]^\lambda \\
&= p(\boldsymbol{\pi}) p(\boldsymbol{\eta}) \times \prod_{k=1}^K \prod_{i=1}^{N^k} p(\mathbf{D}_i^l | Z_i = k, \boldsymbol{\eta}) p(Z_i = k | \boldsymbol{\pi}) \\
&\quad \times \left[ \prod_{i=1}^{N^u} \sum_{k=1}^K \left\{ p(\mathbf{D}_i^u | Z_i = k, \boldsymbol{\eta}) p(Z_i = k | \boldsymbol{\pi}) \right\} \right]^\lambda \\
&\propto \underbrace{\prod_{k=1}^K \left\{ \pi_k^{\alpha_k - 1} \prod_{v=1}^V \eta_{vk}^{\beta_{kv} - 1} \right\}}_{\text{prior}} \times \underbrace{\prod_{k=1}^K \prod_{i=1}^{N^k} \left\{ \prod_{v=1}^V \eta_{vk}^{D_{iv}} \times \pi_k \right\}}_{\text{labeled doc. likelihood}} \times \underbrace{\left[ \prod_{i=1}^{N^u} \sum_{k=1}^K \left\{ \prod_{v=1}^V \eta_{vk}^{D_{iv}} \times \pi_k \right\} \right]^\lambda}_{\text{unlabeled doc. likelihood}}
\end{aligned} \tag{13}$$

The Q function (the expectation of the complete log-likelihood) is

$$\begin{aligned}
Q &\equiv \mathbb{E}_{\mathbf{Z}|\boldsymbol{\pi}^{(t)}, \boldsymbol{\eta}^{(t)}, D, C}[p(\boldsymbol{\pi}, \boldsymbol{\eta}, \mathbf{Z}|\mathbf{D}, \mathbf{C})] \\
&= \sum_{k=1}^K \left[ (\alpha_k - 1) \log \pi_k^{(t)} + \sum_{v=1}^V \left\{ (\beta_{kv} - 1) \log \eta_{vk}^{(t)} \right\} \right] \\
&\quad + \sum_{k=1}^K \sum_{i=1}^{N^k} \left\{ \sum_{v=1}^V D_{iv} \log \eta_{vk}^{(t)} + \log \pi_k^{(t)} \right\} \\
&\quad + \lambda \left[ \sum_{i=1}^{N^u} \sum_{k=1}^K p_{ik} \left\{ \sum_{v=1}^V D_{iv} \log \eta_{vk}^{(t)} + \log \pi_k^{(t)} \right\} \right]
\end{aligned} \tag{14}$$

The posterior probability of  $Z_i = k$ ,  $p_{ik}$ , is

$$p_{ik} = \frac{\prod_{v=1}^V \eta_{vk}^{D_{iv}} \times \pi_k}{\sum_{k=1}^K \left[ \prod_{v=1}^V \eta_{vk}^{D_{iv}} \times \pi_k \right]} \tag{15}$$

M step estimators are The updating equation for  $\pi$  is the following.

$$\hat{\pi}_k \propto \alpha_k - 1 + N^k + \lambda \sum_{i=1}^{N^u} p_{ik} \tag{16}$$

The updating equation for  $\eta$  is the following.

$$\hat{\eta}_{vk} \propto (\beta_k - 1) + \sum_{i=1}^{N^k} D_{iv} + \lambda \sum_{i=1}^{N^u} p_{ik} D_{iv} \tag{17}$$

Note that we downweight the information from the unlabeled documents by  $\lambda$ , to utilize more reliable information from labeled documents.

## C.2 Results




Figure C.1: **Multiclass Classification Results.**

The darker lines show the results with *activeText* and the lighter lines show the results with SVM. The solid lines use active sampling to decide the next set of documents to be labeled, and the dashed lines use random (passive) sampling. The y-axis indicates the out-of-sample F1 score and the x-axis show the number of sampling steps. The left column shows the results on BBC corpus, where the target classes are “Politics,” “Entertainment,” “Business,” “Sports,” and “Technology.” “Politics” class has 5% of the total dataset, and the rest 95% is evenly split across the rest of classes. The right column shows the results on the Supreme Court corpus, where the target classes are “Criminal Procedure” (32.4% of the corpus), “Civil Rights” (21.4%), “Economic Activity” (22.2%), “Judicial Power” (15.4%), “First Amendment” (8.6%).” In our model, we set the number of latent clusters to be the same as the classification categories and linked each latent cluster to one classification category. *activeText* performs the best across the four specifications on both corpora.

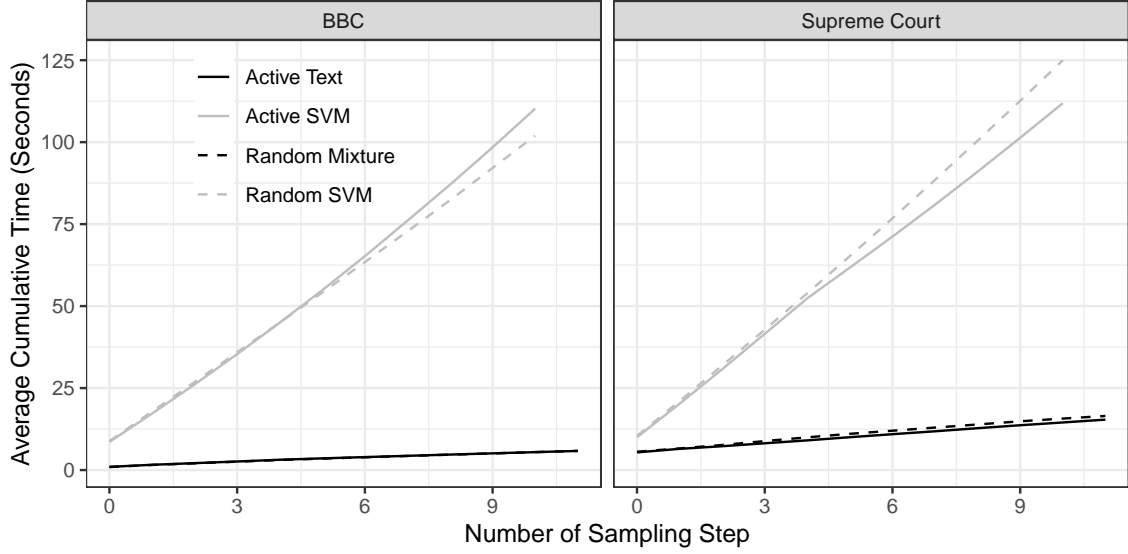


Figure C.2: **Time comparison of Multiclass Classification Results.**

The darker lines show the results with *activeText* and the lighter lines show the results with SVM. The solid lines use active sampling to decide the next set of documents to be labeled, and the dashed lines use random (passive) sampling. The y-axis indicates the average cumulative computational time and the x-axis shows the number of sampling steps. The left column shows the results on BBC corpus, and the right column shows the results on the Supreme Court corpus. *activeText* is much faster than SVM in multiclass classification. This is because multiclass classification with SVM requires fitting the model repeatedly at least the same time as the number of target classes. By contrast, *activeText* requires to fit only once regardless of the number of target classes.

## D Model Specifications and Description of the Datasets in the Validation Performance

We explain our decisions regarding pre-processing steps, model evaluation, and model specifications, followed by a detailed discussion of the results for each dataset.

### D.1 Pre-processing

We employ the same pre-processing step for each of the four datasets using the *R* package *Quanteda*.<sup>1</sup> For each dataset, we construct a *document-feature matrix* (DFM), where each row is a document and each column is a feature. Each feature is a stemmed unigram. We remove stopwords, features that occur extremely infrequently, as well as all features under 4 characters.

To generate dataset with the proportion of positive class  $p$  (e.g. 5% or 50%), we randomly sample documents from the original dataset so that it achieves the proportion of the positive class  $p$ . Suppose the number of documents in the original dataset is  $N$  with  $N_{pos}$  and  $N_{neg}$  the number of positive and negative documents, respectively. We compute  $M_{pos} = \text{floor}(Np)$  and  $M_{neg} = N - M_{pos}$  as the ideal numbers of positive and negative documents. While  $M_{pos} > N_{pos}$  or  $M_{neg} > N_{neg}$ , we decrement  $M_{pos}$  and  $M_{neg}$  keeping the positive proportion to  $p$ . With  $M_{pos} < N_{pos}$  and  $M_{neg} < N_{neg}$ , we sample  $M_{pos}$  positive documents and  $M_{neg}$  negative documents from the original dataset. Finally, combine the sampled positive and negative documents to obtain the final dataset.

### D.2 Datasets

**BBC News** The BBC News Dataset is a collection of 2,225 documents from 2004 to 2005 available at the BBC news website (Greene and Cunningham, 2006). This dataset is divided equally into five topics: business, entertainment, politics, sport, and technology. The classification exercise is to correctly predict whether or not an article belongs to the ‘politics’ topic.

**Wikipedia Toxic Comments** The Wikipedia Toxic Comments dataset is a dataset made up of conversations between Wikipedia editors in Wikipedia’s internal forums. The dataset was made openly available as part of a Kaggle competition,<sup>2</sup> and was used as a principle dataset of investigation by Miller et al. (2020). The basic classification task is to label a given speech as toxic or not, where toxicity is defined as including harassment and/or abuse

---

<sup>1</sup>See <https://quanteda.io>

<sup>2</sup>See <https://www.kaggle.com/c/jigsaw-toxic-comment-classification-challenge>

of other users.<sup>3</sup> The complete dataset is comprised of roughly 560,000 documents, roughly 10 percent of which are labeled as toxic.

**Supreme Court Cases** The Supreme Court Rulings dataset is a collection of the text of 2000 US Supreme Court rulings between 1946 and 2012. We use the majority opinion of each case and the text was obtained through Caselaw Access Project.<sup>4</sup> For the classification label, we use the categories created by the Supreme Court Database.<sup>5</sup> The classification exercise here is to correctly identify rulings that are categorized as ‘criminal procedure’, which is the largest category in the corpus (26% of all rulings).

**Human Rights Allegation** Human Rights Allegation dataset contains more than 2 million sentences of human rights reports in 196 countries between 1996 and 2016, produced by Amnesty International, Human Rights Watch and the US State Department (Cordell et al., 2021). The classification goal is to identify sentences with physical integrity rights allegation (16% of all reports). Example violations of physical integrity rights include torture, extrajudicial killing, and arbitrary arrest and imprisonment.

---

<sup>3</sup>While the dataset also contains finer gradation of ‘types’ of toxicity, we like Miller et al. (2020) stick to the binary toxic-or-not classification task.

<sup>4</sup><https://case.law>

<sup>5</sup>For a full list of categories, see <http://www.supremecourtdatabase.org/documentation.php?var=issueArea>.

## E Additional Results on Classification Performance

To complement the results presented in Figure 1 in the main text, Table E.1 presents the results (across datasets) of fitting our model at the initial (iteration 0) and last active step (iteration 30). It is clear from the table that the improvements *activeText* brings in terms of the F1-score, precision, and recall. Furthermore, after labeling 600 documents (20 per iteration), uncertainty sampling outperforms random sampling across evaluation metrics, which empirically validates the promise of active learning in terms of text classification.

Table E.1: **Classification Performance: Uncertainty vs Random Sampling with  $\lambda = 0.001$**

Dataset	Active Step	Uncertainty Sampling			Random Sampling		
		Precision	Recall	F1-score	Precision	Recall	F1-score
Wikipedia	0	0.71	0.13	0.22	0.71	0.13	0.22
	30	0.71	0.54	0.61	0.45	0.56	0.50
BBC	0	0.33	0.86	0.48	0.33	0.86	0.48
	30	0.92	0.96	0.94	0.92	0.94	0.93
Supreme Court	0	0.46	0.98	0.63	0.46	0.98	0.63
	30	0.85	0.91	0.88	0.75	0.96	0.84
Human Rights	0	0.61	0.01	0.02	0.61	0.01	0.02
	30	0.53	0.42	0.47	0.46	0.44	0.45

Similarly, and as noted in the main text, our results appear to be not too sensitive to the selection of the weighting parameter  $\lambda$ , provided that its value remains small. Figures E.1 confirms this finding. After 30 active steps, the performance of *activeText* is better in terms of F1-score when  $\lambda = 0.001$  if compared to  $\lambda = 0.01$






Figure E.1: **Classification Results with 2 Clusters and  $\lambda = 0.01$  vs  $\lambda = 0.001$ .** The darker lines show the results with  $\lambda = 0.001$  and the lighter lines show  $\lambda = 0.01$ . The columns correspond to various proportion of positive labels in the corpus. The y-axis indicates the out-of-sample F1 score and the x-axis show the number of sampling steps. The smaller the value of  $\lambda$  the better the performance of our model.

## F Main Results when Varying Positive Class Rate

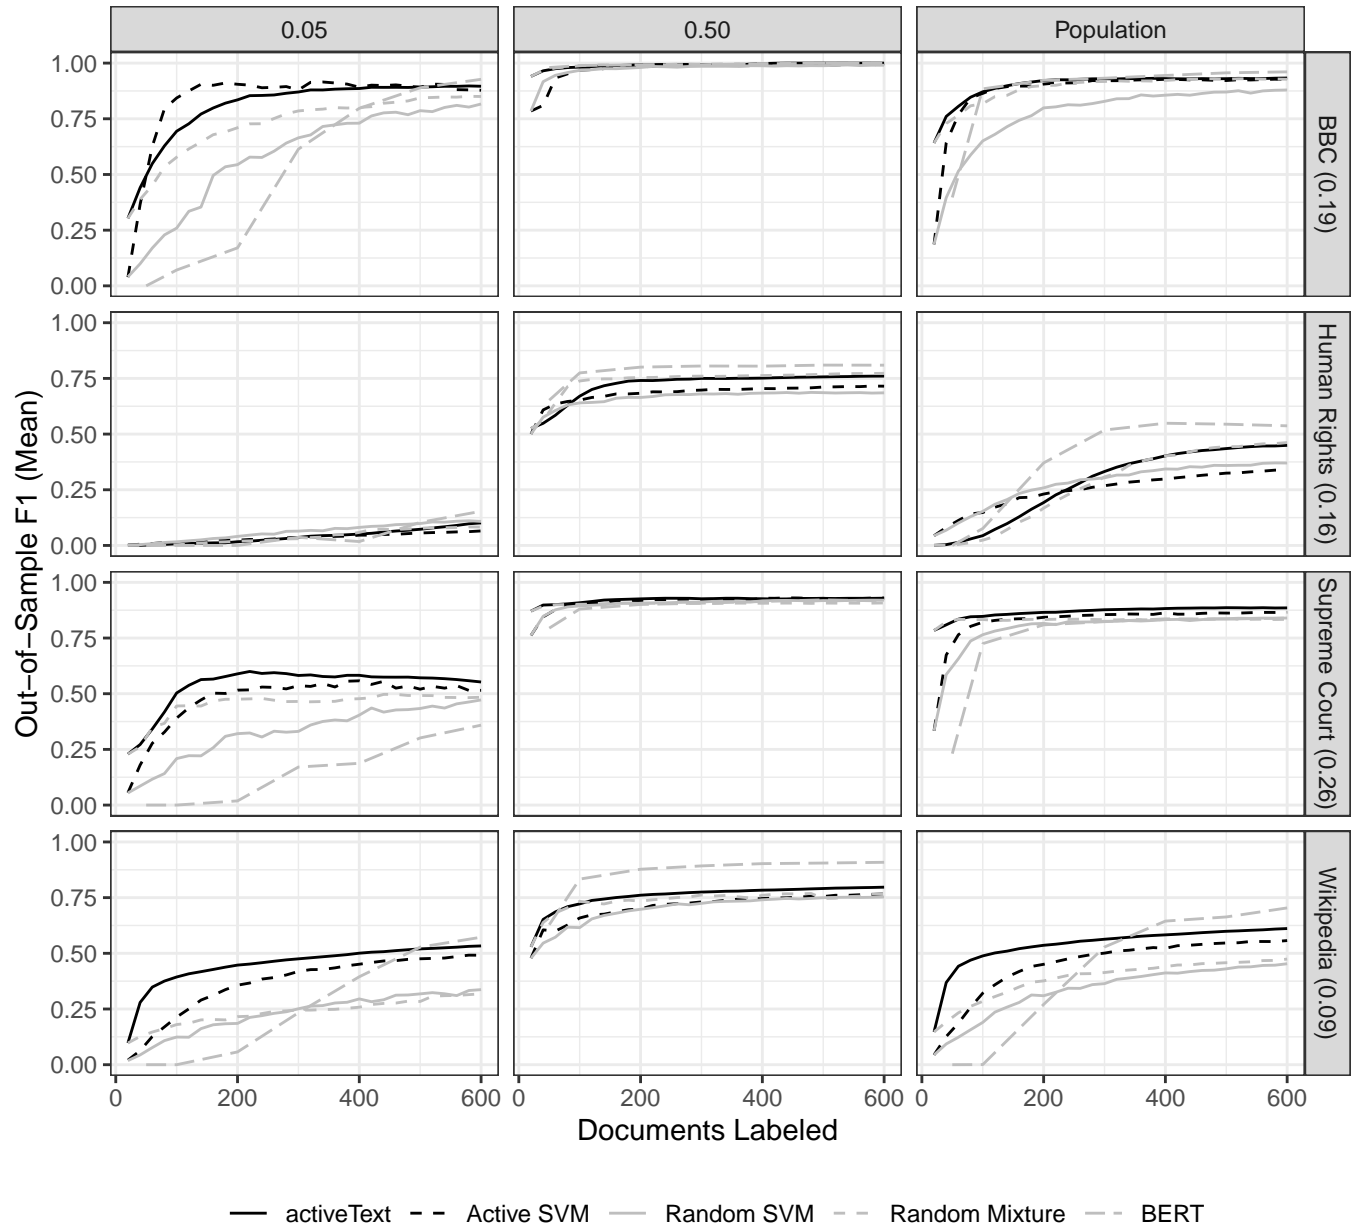


Figure F.1: Replication of F1 performance from Figures 2 and 3 with 0.05, 0.5, and population positive class rate

## G Visual Demonstration of Active Keyword

Figure G.1 illustrates how the word-class matrix  $\boldsymbol{\eta}$  is updated with and without keywords across iterations. A subset of the keywords supplied is labeled and highlighted by black dots. The x-axis shows the log of  $\eta_{v1}/\eta_{v0}$ , where  $\eta_{v1}$  corresponds the probability of observing the word  $v$  in a document with a positive label and  $\eta_{v0}$  for a document with a negative label. The high value in the x-axis means that a word is more strongly associated with positive labels. The y-axis is the log of word frequency. A word with high word frequency has more influence in shifting the label probability. In the generative model for *activeText*, words that appear often and whose ratio of  $\eta_{vk^*}$  vs  $\eta_{vk}$  is high play a central role in the label prediction. By shifting the value of  $\boldsymbol{\eta}$  of those keywords, we can accelerate the estimation of  $\boldsymbol{\eta}$  and improve the classification performance.




Figure G.1: Update of the Word-class Matrix ( $\eta$ ) with and without Keywords

## H Classification Performance with Mislabels

### H.1 Mislabeled Keywords

The rows correspond to different datasets and the columns correspond to various values of  $\gamma$ , which controls the degree of keyword upweighting. The y-axis indicates the out-of-sample F1 score and the x-axis shows the number of sampling steps. At each sampling step, 20 documents are labeled. We use  $\lambda = 0.001$  to downweight information from unlabeled documents. The lines correspond to different levels of mislabels at the keyword labeling. At each iteration, 10 candidate keywords are proposed, and a hypothetical oracle decides if they are indeed keywords or not. ‘True’ keywords are defined in the same way as in Section 4.3. In other words, a candidate keyword  $v$  for the positive class is a ‘true’ keyword, if the value of  $\eta_{v,k}/\eta_{v,k'}$  is above 90% quantile, where  $k$  is the positive class and  $k'$  is the negative class, and this  $\eta$  is what we obtain by training the model with the full labels. The same goes for the negative class. When the probability of mislabeling keywords is  $p\%$ , an oracle makes a mistake in the labeling with probability  $p$ . Specifically, if a candidate keyword  $v$  is a ‘true’ keyword, the oracle would not label  $v$  as a keyword with probability  $p$ . Likewise, if a candidate keyword  $v$  is not a ‘true’ keyword, they would label  $v$  as a keyword.




Figure H.1: Classification Results with Mislabels in Active Keywords

## H.2 Mislabeled Documents

In this section, we present results about the effect of ‘honest’ (random) mislabeling of documents on the mapping of documents to classes. As Figure H.2 shows, as the proportion of mislabels increases, the classification performance of *activeText* decreases.

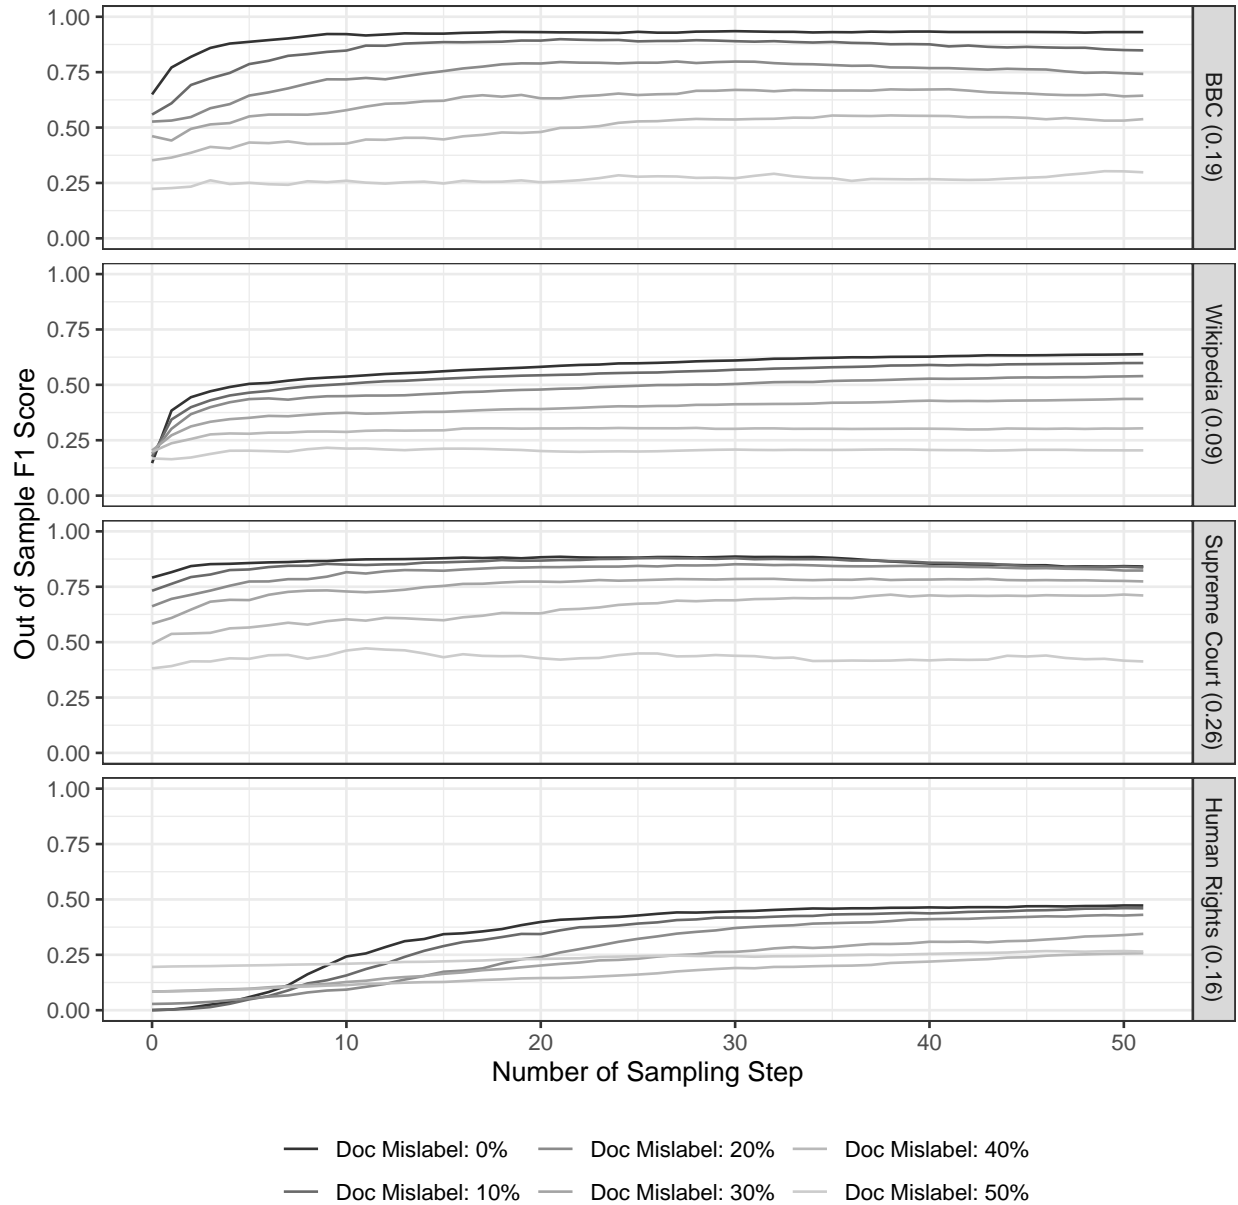


Figure H.2: **Classification Results with Mislabeled in Active Document Labeling**  
The rows correspond to different datasets. The y-axis indicates the out-of-sample F1 score and the x-axis shows the number of sampling steps. 20 documents are labeled at each sampling step. The colors correspond to different levels of mislabels in the labeling of documents. We find that as the proportion of mislabels increases, the classification performance of *activeText* decreases.



# I Comparison of the predictions between *activeText* and xgboost predictions for the Gohdes (2020) data

Table I.1 shows the confusion matrix between the prediction based on *activeText* and the prediction by xgboost used in the original paper. Most observations fall in the diagonal cells of the matrix, and the correlation between the two predictions is quite high (0.93). One difference is that *activeText* classifies more documents to target killings compared to the original predictions. Note that either prediction claims the ground truth. Both are the results of different classifiers.

		Original		
		untargeted	targeted	non-government
<i>activeText</i>	untargeted	50327	411	135
	targeted	1630	10044	31
	non-government	382	34	2280

Table I.1: Confusion matrix between *activeText* and xgboost predictions




Figure I.1: **Scatter plot of the dependent variable between the one constructed by *activeText* vs. the original** The author performs a binomial logit regression where the dependent variable is the ratio of the number of targeted killings to the total number of government killings. We compare the dependent variable used in the original paper vs. the one we constructed using *activeText* . The 45-degree line (in red) corresponds to equality between measures. We can see that most observations lie around the 45-degree line while there are some values in the upper triangle. This suggests that *activeText* yields a similar dependent variable to the original one, while there may be some overestimations of the proportion of target killing with *activeText* .

## J Regression Table in Gohdes (2020)

Table J.1 is the original regression table reported in Gohdes (2020) while Table J.2 is a replication of the original table using *activeText* . In both tables, the coefficients on the Internet access variable are positive and statistically significant, which match the author’s substantive conclusion. One may wonder why the absolute values of the coefficients on the IS and Internet is larger in Table J.2. However, we believe that this is because the number of observations in the IS control is small (51) and there is almost no variation of the Internet access variable within the observations with IS control, as shown in Figure J.1.

	I	II	III	IV	V	VI	VII
Intercept	-2.340*** (0.205)	-2.500*** (0.267)	-0.899* (0.403)	-0.410 (0.521)	-0.019 (0.357)	-1.308 (1.057)	-3.013** (1.103)
Internet access (3G)	0.224* (0.095)	0.231* (0.094)	0.200* (0.085)	0.205* (0.087)	0.265* (0.113)	0.313** (0.116)	0.909*** (0.124)
% Govt control							0.016*** (0.004)
Internet (3G) * % Govt control							-0.014*** (0.001)
Govt control	0.774* (0.332)	0.803** (0.272)	1.167*** (0.284)	1.180*** (0.288)	0.080 (0.344)	0.856** (0.313)	0.811*** (0.237)
IS control	2.027*** (0.435)	1.644*** (0.462)	1.045* (0.421)	-0.324 (0.209)	0.432 (0.414)	0.787 (0.418)	-0.663** (0.221)
Kurd control	0.386 (0.594)	-0.243 (0.843)	-0.506 (0.760)	-1.331 (1.134)	-0.402 (0.745)	0.033 (0.802)	-0.616 (0.432)
Opp control	1.160*** (0.298)	1.252*** (0.317)	0.727* (0.293)	0.759* (0.296)	-0.700* (0.283)	-0.281 (0.342)	-0.176 (0.164)
Internet (3G) * Govt control	-0.163 (0.132)	-0.182 (0.117)	-0.327** (0.119)	-0.324** (0.122)	-0.104 (0.133)	-0.358** (0.120)	
Internet (3G) * IS control	-1.798*** (0.220)	-1.525*** (0.281)	-1.377*** (0.251)		-1.391*** (0.264)	-1.336*** (0.261)	
Internet (3G) * Kurd control	-0.133 (0.444)	0.336 (0.649)	0.093 (0.569)	0.895 (0.936)	-0.052 (0.553)	-0.202 (0.527)	
Internet (3G) * Opp. control	-0.605*** (0.159)	-0.722*** (0.173)	-0.511** (0.157)	-0.533*** (0.158)	0.316* (0.151)	0.286 (0.186)	
# Killings (log)			-0.273*** (0.054)	-0.271*** (0.055)	-0.354*** (0.051)	-0.412*** (0.072)	-0.584*** (0.074)
Govt gains				0.643 (0.385)			
Govt losses				0.632 (0.413)			
Christian					0.068 (0.111)	0.345** (0.116)	0.398*** (0.110)
Alawi					1.479** (0.522)	-1.167*** (0.177)	-0.812*** (0.176)
Druze					-0.634*** (0.191)	-0.302 (0.191)	0.135 (0.190)
Kurd					-0.659*** (0.194)	-0.542* (0.237)	-0.580** (0.212)
Internet (3G) * Alawi					-0.909*** (0.163)		
Pop (log)						0.196 (0.149)	0.408** (0.150)
Unempl. (%)						-0.016 (0.012)	-0.002 (0.012)
AIC	11956.847	9993.704	9665.749	9495.591	7671.979	7873.915	7327.796
BIC	12001.524	10239.427	9915.941	9744.552	7944.509	8150.913	7595.858
Log Likelihood	-5968.424	-4941.852	-4776.875	-4691.796	-3774.990	-3874.958	-3603.898
Deviance	9519.651	7466.508	7136.554	7026.891	5132.784	5332.720	4790.601
Num. obs.	640	640	640	626	640	640	640

\*\*\* $p < 0.001$ ; \*\* $p < 0.01$ ; \* $p < 0.05$ . Reference category: Contested control. Governorate-clustered SEs.

Table J.1: Table 1 in Gohdes 2020: Original table

	I	II	III	IV	V	VI	VII
Intercept	-2.196*** (0.197)	-2.428*** (0.242)	-0.795* (0.390)	-0.351 (0.490)	-0.037 (0.348)	-1.141 (1.229)	-2.695* (1.227)
Internet access (3G)	0.277** (0.091)	0.282*** (0.081)	0.242** (0.075)	0.250** (0.077)	0.342*** (0.103)	0.369*** (0.107)	0.853*** (0.118)
% Govt control							0.015*** (0.004)
Internet (3G) * % Govt control							-0.013*** (0.001)
Govt control	0.625* (0.319)	0.672** (0.255)	1.048*** (0.269)	1.058*** (0.273)	0.151 (0.358)	0.843** (0.300)	0.559* (0.249)
IS control	15.157*** (1.123)	15.688*** (1.148)	15.072*** (1.136)	-0.275 (0.200)	14.551*** (1.132)	14.877*** (1.134)	-0.600** (0.209)
Kurd control	0.795 (0.516)	0.099 (0.729)	-0.227 (0.671)	-0.440 (1.119)	-0.157 (0.677)	0.334 (0.744)	-0.369 (0.405)
Opp control	0.978*** (0.294)	1.134*** (0.304)	0.594* (0.284)	0.634* (0.289)	-0.606* (0.270)	-0.197 (0.322)	-0.278 (0.155)
Internet (3G) * Govt control	-0.169 (0.126)	-0.190 (0.103)	-0.334** (0.108)	-0.335** (0.111)	-0.183 (0.131)	-0.408*** (0.111)	
Internet (3G) * IS control	-14.829*** (1.080)	-15.506*** (1.096)	-15.351*** (1.090)		-15.392*** (1.091)	-15.330*** (1.091)	
Internet (3G) * Kurd control	-0.400 (0.324)	0.138 (0.514)	-0.080 (0.463)	0.134 (0.940)	-0.240 (0.473)	-0.366 (0.460)	
Internet (3G) * Opp. control	-0.542*** (0.159)	-0.688*** (0.164)	-0.468** (0.150)	-0.497** (0.152)	0.181 (0.145)	0.149 (0.176)	
# Killings (log)			-0.278*** (0.053)	-0.274*** (0.054)	-0.356*** (0.051)	-0.415*** (0.071)	-0.567*** (0.073)
Govt gains				0.512 (0.349)			
Govt losses				0.730* (0.334)			
Christian					0.092 (0.115)	0.352** (0.113)	0.369*** (0.105)
Alawi					1.329* (0.528)	-0.928*** (0.167)	-0.585*** (0.168)
Druze					-0.628** (0.196)	-0.310 (0.197)	0.063 (0.209)
Kurd					-0.565** (0.204)	-0.502* (0.227)	-0.615** (0.207)
Internet (3G) * Alawi					-0.782*** (0.164)		
Pop (log)						0.185 (0.167)	0.391* (0.168)
Unempl. (%)						-0.019 (0.012)	-0.007 (0.012)
AIC	12050.644	10116.531	9739.975	9570.556	8038.596	8197.433	7735.527
BIC	12095.321	10362.255	9990.166	9819.517	8311.125	8474.431	8003.589
Log Likelihood	-6015.322	-5003.266	-4813.988	-4729.278	-3958.298	-4036.717	-3807.763
Deviance	9500.059	7475.946	7097.391	6986.658	5386.011	5542.849	5084.942
Num. obs.	640	640	640	626	640	640	640

\*\*\* $p < 0.001$ ; \*\* $p < 0.01$ ; \* $p < 0.05$ . Reference category: Contested control. Governorate-clustered SEs.

Table J.2: Table 1 in Gohdes 2020: Reanalysis with *activeText*




Figure J.1: **Histogram of the Internet (3G) variable by the IS control in the original data** The left histogram is the distribution of the Internet (3G) variable for the observation under IS control, and the right one is not under IS control. The number of observations with IS control is only 51 out of the total observation of 640. In addition, among those with IS control, all observations except one takes the same value for the Internet access variable. This suggests that the regression coefficient on the interaction of IS control and Internet access can be highly unstable.

## K Effect of Labeling More Sentences for the Park et al. (2020) Reanalysis

In this section, we present additional results mentioned in the main text about our reanalysis of Park et al. (2020).




Figure K.1: Using the Difference in Out of Sample F1 Score to Decide a Stopping Point.




Figure K.2: Replication of Figure 1 in Park et al. (2020): The Relationship Between Information Density and Average Sentiment Score Across Different Settings for the Total Number of Labeled Documents.

## References

- Cordell, R., Clay, K. C., Fariss, C. J., Wood, R. M., and Wright, T. (2021), “Recording repression: Identifying physical integrity rights allegations in annual country human rights reports,” *International Studies Quarterly*, .
- Dempster, A. P., Laird, N. M., and Rubin, D. B. (1977), “Maximum likelihood from incomplete data via the EM algorithm,” *Journal of the Royal Statistical Society: Series B (Methodological)*, 39(1), 1–22.
- Gohdes, A. R. (2020), “Repression technology: Internet accessibility and state violence,” *American Journal of Political Science*, 64(3), 488–503.
- Greene, D., and Cunningham, P. (2006), Practical Solutions to the Problem of Diagonal Dominance in Kernel Document Clustering,, in *Proc. 23rd International Conference on Machine learning (ICML’06)*, ACM Press, pp. 377–384.
- Larsen, M. D., and Rubin, D. B. (2001), “Iterative Automated Record Linkage Using Mixture Models.,” *Journal of the American Statistical Association.*, 96(453), 32–41.
- Miller, B., Linder, F., and Mebane, W. R. (2020), “Active Learning Approaches for Labeling Text: Review and Assessment of the Performance of Active Learning Approaches,” *Political Analysis*, pp. 1–20.
- Park, B., Greene, K., and Colaresi, M. (2020), “Human Rights are (Increasingly) Plural: Learning the Changing Taxonomy of Human Rights from Large-scale Text Reveals Information Effects,” *American Political Science Review*, 114(3), 888–910.